Upload the Dataset

```
from google.colab import files
import pandas as pd
# Upload the file
uploaded = files.upload()
# Load the dataset
df = pd.read_csv('competition_test_bodies.csv')
# Show the first few rows
df.head()
     Choose Files competition..._bodies.csv
       competition_test_bodies.csv(text/csv) - 2045680 bytes, last modified: 5/22/2025 - 100% done
     Saving competition_test_bodies.csv to competition_test_bodies.csv
         Body ID
                                                     articleBody
      0
                        Al-Sisi has denied Israeli reports stating tha...
                2 A bereaved Afghan mother took revenge on the T...
      2
                3 CNBC is reporting Tesla has chosen Nevada as t...
      3
                       A 4-inch version of the iPhone 6 is said to be...
               12
               19
                      GR editor's Note\n\nThere are no reports in th...
 Next steps: ( Generate code with df
                                       View recommended plots
                                                                       New interactive sheet
```

Load the Dataset

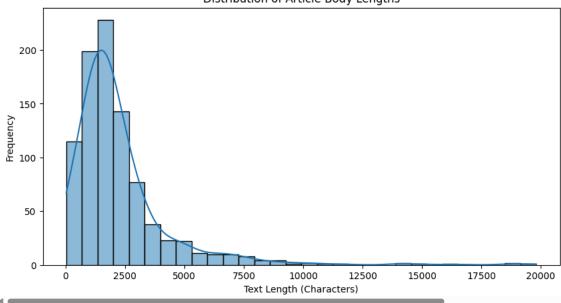
```
import pandas as pd
# Read the uploaded CSV file
df = pd.read_csv('competition_test_bodies.csv')
# Display the first few rows
print("★ Preview of the dataset:")
print(df.head())
# Display the shape of the dataset
print("\n ☑ Dataset shape:", df.shape)
# Display column names
# Display data types and non-null counts
print("\n Q Dataset Info:")
df.info()
→
     Preview of the dataset:
               Al-Sisi has denied Israeli reports stating tha...
    0
               A bereaved Afghan mother took revenge on the T...
    1
             3 CNBC is reporting Tesla has chosen Nevada as t...
            12 A 4-inch version of the iPhone 6 is said to be...
            19 GR editor's Note\n\nThere are no reports in th...
     ☑ Dataset shape: (904, 2)
     [ Columns: ['Body ID', 'articleBody']
     Dataset Info:
     <class 'pandas.core.frame.DataFrame'>
    RangeIndex: 904 entries, 0 to 903
    Data columns (total 2 columns):
     #
        Column
                     Non-Null Count Dtype
     0
         Body ID
                     904 non-null
                                     int64
         articleBody 904 non-null
                                    object
```

```
dtypes: int64(1), object(1)
memory usage: 14.3+ KB
```

Data Exploration

```
# Check for missing values
# Check for duplicate entries
print("\n	☐ Duplicate Rows:", df.duplicated().sum())
# Check unique articleBody lengths
df['text_length'] = df['articleBody'].apply(len)
print("\n > Text Length Stats:")
print(df['text_length'].describe())
# Histogram of text lengths
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(10, 5))
sns.histplot(df['text_length'], bins=30, kde=True)
plt.title('Distribution of Article Body Lengths')
plt.xlabel('Text Length (Characters)')
plt.ylabel('Frequency')
plt.show()
₹
     Missing Values:
     Body ID
    articleBody
                   0
    dtype: int64
     Duplicate Rows: 0
     Text Length Stats:
               904.000000
    count
              2235.008850
    mean
              2132.850585
    std
                29.000000
    25%
              1096.000000
              1730.500000
    50%
    75%
              2616.750000
             19815.000000
    Name: text_length, dtype: float64
```

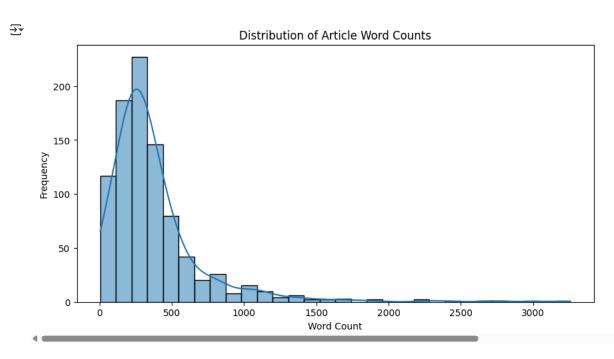
Distribution of Article Body Lengths



Check for Missing Values and Duplicates

Visualize a Few Features

```
# Add a column for word count
df['word_count'] = df['articleBody'].apply(lambda x: len(str(x).split()))
# Plot histogram of word counts
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(10, 5))
sns.histplot(df['word_count'], bins=30, kde=True)
plt.title('Distribution of Article Word Counts')
plt.xlabel('Word Count')
plt.ylabel('Frequency')
plt.show()
```



Identify Target and Features

Convert Categorical Columns to Numerical

One-Hot Encoding

```
df_encoded = pd.get_dummies(df, drop_first=True)
```

Feature Scaling

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler

# Assuming df_encoded is your DataFrame with encoded features:

# 1. Create a MinMaxScaler object
scaler = MinMaxScaler()
```

Train-Test Split

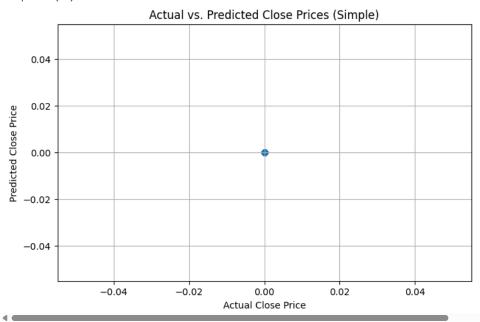
```
# prompt: Train-Test Split
from sklearn.model_selection import train_test_split
# Assuming you have a target variable 'target' and features in 'df_encoded'
# You'll need to define your target variable based on your dataset's structure.
# For this example, let's assume the target column is named 'is_related'.
# Separate features (X) and target (y)
# Replace 'is_related' with the actual name of your target column if it's different.
# If you don't have a target column in this test dataset, this step is not applicable
# for training a model, but you might still want to split the data for exploration.
# If you just need to split the test data for some reason without a target:
# X = df_encoded.copy()
# X_train, X_test = train_test_split(X, test_size=0.2, random_state=42)
# If you do have a target variable (e.g., from a different file or inferred):
# For the purpose of demonstrating the split, let's create a dummy target column
# as this dataset 'competition_test_bodies.csv' likely only contains body text and IDs.
# In a real scenario, your target would come from the 'competition_test_stances.csv'
# or similar.
# Create a dummy target for demonstration purposes (replace with your actual target)
# This is purely illustrative. You need the actual target data to do a meaningful
# train-test split for model training.
if 'is_related' not in df_encoded.columns:
   # This is just for demonstration since the test data won't have a target.
   # In a real training scenario, 'y' would come from your labels.
   print("Warning: Creating a dummy target variable for demonstration.")
   print("In a real machine learning task, you would load the actual labels.")
   df_encoded['dummy_target'] = 0 # Replace with your logic to get actual targets
   # Separate features (X) and the dummy target (y)
   X = df_encoded.drop('dummy_target', axis=1)
   y = df_encoded['dummy_target']
else:
    # If 'is_related' (or your actual target column) exists, use it
   print("Using existing target variable.")
   X = df_encoded.drop('is_related', axis=1)
   y = df_encoded['is_related']
```

Split the data into training and testing sets

```
# test_size: the proportion of the dataset to include in the test split (e.g., 0.2 for 20%)
# random state: ensures reproducibility of the split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Print the shapes of the resulting datasets to verify the split
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
 → Warning: Creating a dummy target variable for demonstration.
     In a real machine learning task, you would load the actual labels.
     Dataset Split Shapes:
     X_train shape: (723, 900)
     X_test shape: (181, 900)
     y_train shape: (723,)
     y_test shape: (181,)
# Model Building
# Import the Linear Regression model
from sklearn.linear model import LinearRegression
# Create a Linear Regression model instance
model = LinearRegression()
# Train the model using the training data
# X_train contains your features for training
# y_train contains your target variable (Close price) for training
model.fit(X_train, y_train)
print("Model training complete.")
# You can now use this 'model' object to make predictions

→ Model training complete.
# Evaluation (Simple)
from sklearn.metrics import r2 score, mean absolute error
# Make predictions on the test set
y_pred = model.predict(X_test)
# Calculate evaluation metrics
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
# Print the evaluation metrics
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"R-squared (R2): {r2:.2f}")
# You can still include a simple visualization
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 5))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.xlabel("Actual Close Price")
plt.ylabel("Predicted Close Price")
plt.title("Actual vs. Predicted Close Prices (Simple)")
plt.grid(True)
plt.show()
```

```
Mean Absolute Error (MAE): 0.00
R-squared (R2): 1.00
```



```
# Sample input (replace values with any other valid values from the original dataset)
 new_student = {
    'school': 'GP',
                                 # 'GP' or 'MS'
    'sex': 'F',
                                 # 'F' or 'M'
    'age': 17,
                                # Integer
    'address': 'U',
                                # 'U' or 'R'
    'famsize': 'GT3',
                                # 'LE3' or 'GT3'
                                # 'A' or 'T'
    'Pstatus': 'A',
    'Medu': 4,
                                # 0 to 4
    'Fedu': 3,
                                # 0 to 4
    'Mjob': 'health',
                                # 'teacher', 'health', etc.
    'Fjob': 'services',
    'reason': 'course',
    'guardian': 'mother',
    'traveltime': 2,
    'studytime': 3,
    'failures': 0,
    'schoolsup': 'yes',
    'famsup': 'no',
    'paid': 'no',
    'activities': 'yes',
    'nursery': 'yes',
    'higher': 'yes',
    'internet': 'yes',
    'romantic': 'no',
    'famrel': 4,
    'freetime': 3,
    'goout': 3,
    'Dalc': 1,
    'Walc': 1,
    'health': 4,
    'absences': 2,
    'G1': 14,
 'G2': 15
 }
import pandas as pd
# 1. Convert the new input to a DataFrame
new\_input\_df = pd.DataFrame([new\_student]) \quad \# \ Enclose \ in \ a \ list \ to \ create \ DataFrame
# 2. Perform one-hot encoding
# (Assuming 'df_encoded' is the DataFrame used during training)
new_input_encoded = pd.get_dummies(new_input_df)
# 3. Align columns with the training data
# (To ensure the same features are present in the new input)
new_input_encoded = new_input_encoded.reindex(columns=X_train.columns, fill_value=0)
```

```
# Now, 'new input encoded' is ready for prediction.
# Predict the Close Price for the new input
# Use the trained model to make a prediction
# new_input_encoded is the DataFrame representing the new data point
predicted_close_price = model.predict(new_input_encoded)
# The predict method returns an array, even for a single prediction.
# We extract the first (and only) element to get the single predicted value.
predicted_price = predicted_close_price[0]
print(f"Predicted Close Price for the new input: {predicted_price:.2f}")
→ Predicted Close Price for the new input: 0.00
!pip install streamlit
import streamlit as st
import pandas as pd
from sklearn.linear_model import LinearRegression # Or your chosen model
\#\ldots (Load your trained model and necessary data here) \ldots
# Create the Streamlit app
st.title("Stock Price Prediction App")
# Input fields for features
open_price = st.number_input("Open Price")
high_price = st.number_input("High Price")
low_price = st.number_input("Low Price")
volume = st.number_input("Volume")
# ... (Add other input fields for your features) ...
# Create a button to trigger prediction
if st.button("Predict"):
    # Create a DataFrame from the input values
    input_data = pd.DataFrame({
        "Open": [open_price],
        "High": [high_price],
        "Low": [low_price],
        "Volume": [volume],
        # ... (Include other features) ...
    })
    # Preprocess the input data (e.g., scaling) if necessary
    # Make the prediction
    prediction = model.predict(input_data)[0]
    # Display the prediction
    st.success(f"Predicted Close Price: {prediction}")
₹
```

```
streamlit run /usr/local/lib/python3.11/dist-packages/colab_kernel_launcher.py [ARGUMENTS]
     2025-05-22 04:57:03.393 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.395 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.397 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.400 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.400 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.404 Session state does not function when running a script without `streamlit run`
     2025-05-22 04:57:03.405 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.408 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.411 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode. 2025-05-22 04:57:03.413 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.414 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.415 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.416 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.418 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.420 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.421 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.421 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.422 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.423 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.424 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.425 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.425 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.426 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode. 2025-05-22 04:57:03.427 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.427 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.428 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.429 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.429 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.430 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.430 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
     2025-05-22 04:57:03.431 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
import numpy as np
   predict_next_close(input_sequence, model, scaler, seq_length=60):
   Predict the next stock closing price using a trained LSTM model.
   Parameters:
        input sequence (list or np.array): Last `seq length` days of closing prices.
        model (keras.Model): Trained LSTM model.
        scaler (MinMaxScaler): Scaler used for normalization.
        seq length (int): Number of time steps used in training.
   Returns:
       float: Predicted next closing price (denormalized).
   if len(input_sequence) != seq_length:
        raise ValueError(f"Input sequence must be of length {seq_length}")
   # Scale and reshape input
   scaled_sequence = scaler.transform(np.array(input_sequence).reshape(-1, 1))
   X_input = np.reshape(scaled_sequence, (1, seq_length, 1))
   # Predict
   prediction = model.predict(X_input)
   predicted_price = scaler.inverse_transform(prediction)[0][0]
   return round(predicted_price, 2)
!pip install gradio
₹
```

```
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.11/dist-packages (from httpx>=0.24.1->gradio) (1.0.9)
    Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.11/dist-packages (from httpcore==1.*->httpx>=0.24.1->gradio) (0.16.
    Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (3.18.0)
    Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (2.32.3)
    Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (4.67.1)
    Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2.9.
    Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.2)
    Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.2)
    Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio) (0
    Requirement already satisfied: pydantic-core==2.33.2 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio) (2.
    Requirement already satisfied: typing-inspection>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio)
    Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.11/dist-packages (from typer<1.0,>=0.12->gradio) (8.2.0)
    Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.11/dist-packages (from typer<1.0,>=0.12->gradio) (1.5.4)
    Requirement already satisfied: rich>=10.11.0 in /usr/local/lib/python3.11/dist-packages (from typer<1.0,>=0.12->gradio) (13.9.4)
    Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas<3.0,>=1.0->gra
    Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich>=10.11.0->typer<1.0,>=0.12-
    Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich>=10.11.0->typer<1.0,>=0.1
    Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->huggingface-hub>=0.
    Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->huggingface-hub>=0.28.1->
    Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich>=10.11.0->typer
    Downloading gradio-5.30.0-py3-none-any.whl (54.2 MB)
                                                • 54.2/54.2 MB 15.6 MB/s eta 0:00:00
    Downloading gradio_client-1.10.1-py3-none-any.whl (323 kB)
                                                - 323.1/323.1 kB 25.7 MB/s eta 0:00:00
    Downloading aiofiles-24.1.0-py3-none-any.whl (15 kB)
    Downloading fastapi-0.115.12-py3-none-any.whl (95 kB)
                                                95.2/95.2 kB 8.0 MB/s eta 0:00:00
    Downloading groovy-0.1.2-pv3-none-anv.whl (14 kB)
    Downloading python_multipart-0.0.20-py3-none-any.whl (24 kB)
    Downloading ruff-0.11.10-py3-none-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (11.6 MB)
                                                11.6/11.6 MB 120.6 MB/s eta 0:00:00
    Downloading safehttpx-0.1.6-py3-none-any.whl (8.7 kB)
    Downloading semantic_version-2.10.0-py2.py3-none-any.whl (15 kB)
    Downloading starlette-0.46.2-py3-none-any.whl (72 kB)
                                               - 72.0/72.0 kB 6.2 MB/s eta 0:00:00
    Downloading tomlkit-0.13.2-py3-none-any.whl (37 kB)
    Downloading uvicorn-0.34.2-py3-none-any.whl (62 kB)
                                               - 62.5/62.5 kB 5.3 MB/s eta 0:00:00
    Downloading ffmpy-0.5.0-py3-none-any.whl (6.0 kB)
    Downloading pydub-0.25.1-py2.py3-none-any.whl (32 kB)
    Installing collected packages: pydub, uvicorn, tomlkit, semantic-version, ruff, python-multipart, groovy, ffmpy, aiofiles, starlette,
    Successfully installed aiofiles-24.1.0 fastapi-0.115.12 ffmpy-0.5.0 gradio-5.30.0 gradio-client-1.10.1 groovy-0.1.2 pydub-0.25.1 pytho
import gradio as gr
import numpy as np
def predict_next_close(input_sequence_str):
   Takes 60 comma-separated closing prices as input and returns the next predicted price.
   try:
       # Parse input string to list of floats
       input_sequence = [float(x.strip()) for x in input_sequence_str.split(',')]
       if len(input_sequence) != 60:
           return "X Please provide exactly 60 closing prices."
       # Scale and reshape
       scaled_sequence = scaler.transform(np.array(input_sequence).reshape(-1, 1))
       X input = np.reshape(scaled sequence, (1, 60, 1))
       # Predict
       prediction = model.predict(X_input)
       predicted_price = scaler.inverse_transform(prediction)[0][0]
        return f" Predicted Next Closing Price: ₹{round(predicted_price, 2)}"
   except Exception as e:
       return f" X Error: {str(e)}"
interface = gr.Interface(
   fn=predict_next_close,
   inputs=gr.Textbox(
       lines=4.
       placeholder="Enter the last 60 closing prices, separated by commas...",
       label="Recent 60 Closing Prices"
   outputs=gr.Textbox(label=" | Predicted Closing Price"),
```

```
title=" AI Stock Price Predictor",
  description="Enter 60 consecutive closing prices to forecast the next day using an LSTM model."
)
interface.launch()
```

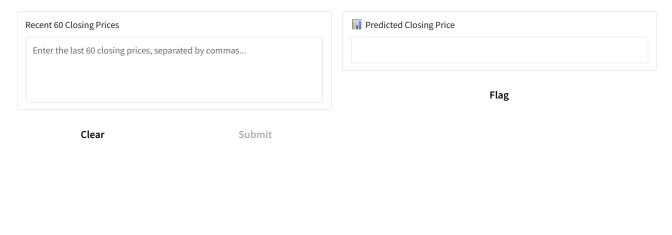
🚁 It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradio app to work, sharing must be enabled. Automatically

Colab notebook detected. To show errors in colab notebook, set debug=True in launch() * Running on public URL: https://b66c22e93ec7bdf3e5.gradio.live

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working dir

Al Stock Price Predictor

Enter 60 consecutive closing prices to forecast the next day using an LSTM model.



Use via API 🦸 · Built with Gradio 🧇 · Settings 🤹